

Only you can see this message



This story's distribution setting is off. [Learn more](#)


# Humor Detection with a BERT Regressor



Theoliao

Apr 25 · 4 min read

Humor (or funniness) detection and generation have always been a challenge in NLP. This could be helpful to build many useful tools. Advising AIs could bring cleverer substitutes for text editing, and text generators might produce more human text. Chatbots and translation AIs might also become more humorous if they “understand” more about humor.



In this article, we try to take it an easy way. That is, to evaluate the level of humor with a regressor trained on some manually graded data. BERT (Bidirectional Encoder Representations from Transformers), as a powerful NLP model with a high-quality pre-trained language model, is used to build such a regressor. Though still far from “understanding” humor, humor can be detected to some extent by a BERT regressor.


## Data

Humicroedit, a dataset introduced in “President Vows to Cut <Taxes> Hair”: Dataset and Analysis of Creative Text Editing for Humorous Headlines, contains collected news headlines with specified words edited to be funnier. Data can be obtained from Headline Humor Dataset — Rochester CS.

The dataset includes a training set and a validation set both with original headlines, edits, grades, and mean grades as shown below. There are 9652 observations in the training set and 2419 in the validation set.

|   | id    | original  | edit     | grades | meanGrade |
|---|-------|---|----------|--------|-----------|
| 0 | 1723  | Thousands of gay and bisexual <men/> convicted... | swans    | 22100  | 1.0       |
| 1 | 12736 | Special <prosecutor/> appointed to Trump Russia   | chef     | 21100  | 0.8       |
| 2 | 12274 | Spanish police detain man and search Ripoll ad... | squad    | 21000  | 0.6       |
| 3 | 8823  | N.Y. Times <reprimands/> reporter for sharing ... | applauds | 32210  | 1.6       |

As we can see, grades here are 5-digit numbers representing 5 manually labeled integer grades from 0 to 3. And what we are interested most is the mean grades. The overall distribution of the mean grades can be visualized as follows.





```

1  # modified from notebook in https://colab.research.google.com/drive/1pTuQhug6Dh19Xa1KB0zUGf4FIc
2
3  from transformers import BertTokenizer
4  import torch
5
6  # Load the BERT tokenizer.
7  tokenizer = BertTokenizer.from_pretrained("bert-base-uncased", do_lower_case=True)
8
9  # Get the lists of sentences and their labels.
10 sentences = train_df['text'].values
11 labels = train_df['meanGrade'].values
12
13 input_ids = []
14 attention_masks = []
15
16 # For every sentence...
17 for sent in sentences:
18     # `encode_plus` will:
19     #   (1) Tokenize the sentence.
20     #   (2) Prepend the `[CLS]` token to the start.
21     #   (3) Append the `[SEP]` token to the end.
22     #   (4) Map tokens to their IDs.
23     #   (5) Pad or truncate the sentence to `max_length`
24     #   (6) Create attention masks for [PAD] tokens.
25     encoded_dict = tokenizer.encode_plus(
26         sent,                                # Sentence to encode.
27         add_special_tokens = True, # Add '[CLS]' and '[SEP]'
28         max_length = 32,                # Pad & truncate all sentences.
29         pad_to_max_length = True,
30         return_attention_mask = True,    # Construct attn. masks.
31         return_tensors = 'pt',         # Return pytorch tensors.
32     )
33
34     # Add the encoded sentence to the list.
35     input_ids.append(encoded_dict['input_ids'])
36
37     # And its attention mask (simply differentiates padding from non-padding).
38     attention_masks.append(encoded_dict['attention_mask'])
39
40 # Convert the lists into tensors.
41 input_ids = torch.cat(input_ids, dim=0)
42 attention_masks = torch.cat(attention_masks, dim=0)
43 labels = torch.tensor(labels)

```

Next, tokenize text in the validation set (and similarly the combined sentence pairs `val_df['text2']`) in the same way.

```
1  # modified from notebook in https://colab.research.google.com/drive/1pTuQhug6Dh19Xa1KB0zUGf4FIc
2
3  sentences_val = val_df['text'].values
4  labels_val = val_df['meanGrade'].values
5
6  # Tokenize all of the sentences and map the tokens to thier word IDs.
7  input_ids_val = []
8  attention_masks_val = []
9
10 # For every sentence...
11 for sent in sentences_val:
12     # `encode_plus` will:
13     #   (1) Tokenize the sentence.
14     #   (2) Prepend the `[CLS]` token to the start.
15     #   (3) Append the `[SEP]` token to the end.
16     #   (4) Map tokens to their IDs.
17     #   (5) Pad or truncate the sentence to `max_length`
18     #   (6) Create attention masks for [PAD] tokens.
19     encoded_dict = tokenizer.encode_plus(
20         sent,                                # Sentence to encode.
21         add_special_tokens = True, # Add '[CLS]' and '[SEP]'
22         max_length = 32,                # Pad & truncate all sentences.
23         pad_to_max_length = True,
24         return_attention_mask = True,    # Construct attn. masks.
25         return_tensors = 'pt',          # Return pytorch tensors.
26     )
27
28     # Add the encoded sentence to the list.
29     input_ids_val.append(encoded_dict['input_ids'])
30
31     # And its attention mask (simply differentiates padding from non-padding).
32     attention_masks_val.append(encoded_dict['attention_mask'])
33
34 # Convert the lists into tensors.
35 input_ids_val = torch.cat(input_ids_val, dim=0)
36 attention_masks_val = torch.cat(attention_masks_val, dim=0)
37 labels_val = torch.tensor(labels_val)
```

Iterators for our dataset are created using the torch DataLoader class. This helps save memory since not all data need to be loaded with an iterator. The batch size is set as 32 here.

```

1  # modified from notebook in https://colab.research.google.com/drive/1pTuQhug6Dh19Xa1KB0zUGf4FIc
2
3  from torch.utils.data import TensorDataset
4  from torch.utils.data import DataLoader, RandomSampler, SequentialSampler
5
6  train_dataset = TensorDataset(input_ids, attention_masks, labels)
7  val_dataset = TensorDataset(input_ids_val, attention_masks_val, labels_val)
8
9  # The DataLoader needs to know our batch size for training, so we specify it
10 # here. For fine-tuning BERT on a specific task, the authors recommend a batch
11 # size of 16 or 32.
12 batch_size = 32
13
14 # Create the DataLoaders for our training and validation sets.
15 # We'll take training samples in random order.
16 train_dataloader = DataLoader(
17     train_dataset, # The training samples.
18     sampler = RandomSampler(train_dataset), # Select batches randomly
19     batch_size = batch_size # Trains with this batch size.
20 )
21
22 # For validation the order doesn't matter, so we'll just read them sequentially.
23 validation_dataloader = DataLoader(
24     val_dataset, # The validation samples.
25     sampler = SequentialSampler(val_dataset), # Pull out batches sequentially.
26     batch_size = batch_size # Evaluate with this batch size.
27 )

```

Next, we need to load the pre-trained BERT model as a regressor. Here we use the *BertForSequenceClassification* class and set the number of labels to be 1, which actually makes it a regressor. The model is set to store double values for a regression task.

```

1  # modified from notebook in https://colab.research.google.com/drive/1pTuQhug6Dh19Xa1KB0zUGf4FIc
2
3  from transformers import BertForSequenceClassification, AdamW, BertConfig
4
5  model = BertForSequenceClassification.from_pretrained(
6      'bert-base-uncased', # Use the 12-layer BERT model, with an uncased vocab.
7      num_labels = 1,
8      output_attentions = False, # Whether the model returns attentions weights.
9      output_hidden_states = False, # Whether the model returns all hidden-states.
10 )
11
12 # Tell pytorch to run this model on the GPU.

```

```

13 model.cuda()
14
15 model = model.double()

```

6 nv hosted with ❤️ by GitHub

[view raw](#)

Parameters are set as follows. Here we use a learning rate of  $2e-5$ , adam epsilon of  $1e-8$ , and 4 epochs.

```

1 # modified from notebook in https://colab.research.google.com/drive/1pTuQhug6Dh19Xa1KB0zUGf4FIc
2
3 from transformers import get_linear_schedule_with_warmup
4
5 optimizer = AdamW(model.parameters(),
6                   lr = 2e-5, # args.learning_rate - default is 5e-5,
7                   eps = 1e-8 # args.adam_epsilon - default is 1e-8.
8                   )
9
10 # Number of training epochs. The BERT authors recommend between 2 and 4.
11 epochs = 4
12
13 # Total number of training steps is [number of batches] x [number of epochs].
14 # (Note that this is not the same as the number of training samples).
15 total_steps = len(train_dataloader) * epochs
16
17 # Create the learning rate scheduler.
18 scheduler = get_linear_schedule_with_warmup(optimizer,
19                                             num_warmup_steps = 0, # Default value in run_glue.py
20                                             num_training_steps = total_steps)

```

7 nv hosted with ❤️ by GitHub

[view raw](#)

Create a function to format time for later use.

```

1 # modified from notebook in https://colab.research.google.com/drive/1pTuQhug6Dh19Xa1KB0zUGf4FIc
2
3 import time
4 import datetime
5
6 def format_time(elapsed):
7     '''
8     Takes a time in seconds and returns a string hh:mm:ss
9     '''
10    # Round to the nearest second.
11    elapsed_rounded = int(round((elapsed)))
12

```

```
13     # Format as hh:mm:ss
14     return str(datetime.timedelta(seconds=elapsed_rounded))
```

8 nv hosted with  by GitHub

[view raw](#)

Finally, we can train the regressor. A specific seed value is set so that the result might be reproduced. After every, the RMSE(Root Mean Squared Error) score on the validation set is evaluated to show the performances.

```
1  # modified from notebook in https://colab.research.google.com/drive/1pTuQhug6Dh19Xa1KB0zUGf4FI
2
3  import random
4  from sklearn.metrics import mean_squared_error
5
6  # This training code is based on the `run_glue.py` script here:
7  # https://github.com/huggingface/transformers/blob/5bfcd0485ece086ebcbed2d008813037968a9e58/ex
8
9  # Set the seed value all over the place to make this reproducible.
10 seed_val = 42
11
12 random.seed(seed_val)
13 np.random.seed(seed_val)
14 torch.manual_seed(seed_val)
15 torch.cuda.manual_seed_all(seed_val)
16
17 # We'll store a number of quantities such as training and validation loss,
18 # validation accuracy, and timings.
19 training_stats = []
20
21 # Measure the total training time for the whole run.
22 total_t0 = time.time()
23
24 # For each epoch...
25 for epoch_i in range(0, epochs):
26
27     # # =====
28     # #           Training
29     # # =====
30
31     # Perform one full pass over the training set.
32
33     print("")
34     print('=====Epoch {:} / {:} ====='.format(epoch_i + 1, epochs))
35     print('Training...')
36
37     # Measure how long the training epoch takes.
```



```

38     t0 = time.time()
39
40     # Reset the total loss for this epoch.
41     total_train_loss = 0
42
43     # Put the model into training mode. Don't be misled--the call to
44     # `train` just changes the *mode*, it doesn't *perform* the training.
45     # `dropout` and `batchnorm` layers behave differently during training
46     # vs. test (source: https://stackoverflow.com/questions/51433378/what-does-model-train-do-
47     model.train()
48
49     # For each batch of training data...
50     for step, batch in enumerate(train_dataloader):
51
52         # Progress update every 40 batches.
53         if step % 40 == 0 and not step == 0:
54             # Calculate elapsed time in minutes.
55             elapsed = format_time(time.time() - t0)
56
57             # Report progress.
58             print(' Batch {:>5,} of {:>5,}. Elapsed: {:.}'.format(step, len(train_dataloader),
59
60             # Unpack this training batch from our dataloader.
61             #
62             # As we unpack the batch, we'll also copy each tensor to the GPU using the
63             # `to` method.
64             #
65             # `batch` contains three pytorch tensors:
66             # [0]: input ids
67             # [1]: attention masks
68             # [2]: labels
69             b_input_ids = batch[0].to(device)
70             b_input_mask = batch[1].to(device)
71             b_labels = batch[2].to(device)
72
73             # Always clear any previously calculated gradients before performing a
74             # backward pass. PyTorch doesn't do this automatically because
75             # accumulating the gradients is "convenient while training RNNs".
76             # (source: https://stackoverflow.com/questions/48001598/why-do-we-need-to-call-zero-grad-
77             model.zero_grad()
78
79             # Perform a forward pass (evaluate the model on this training batch).
80             # The documentation for this `model` function is here:
81             # https://huggingface.co/transformers/v2.2.0/model\_doc/bert.html#transformers.BertForSequenceClassification
82             # It returns different numbers of parameters depending on what arguments
83             # are given and what flags are set. For our usage here, it returns
84             # the loss (because we provided labels) and the "logits"--the model
85             # outputs

```

```

85         # outputs prior to activation.
86         loss, logits = model(b_input_ids,
87                               token_type_ids=None,
88                               attention_mask=b_input_mask,
89                               labels=b_labels)
90
91         # Accumulate the training loss over all of the batches so that we can
92         # calculate the average loss at the end. `loss` is a Tensor containing a
93         # single value; the `.item()` function just returns the Python value
94         # from the tensor.
95         total_train_loss += loss.item()
96
97         # Perform a backward pass to calculate the gradients.
98         loss.backward()
99
100        # Clip the norm of the gradients to 1.0.
101        # This is to help prevent the "exploding gradients" problem.
102        torch.nn.utils.clip_grad_norm_(model.parameters(), 1.0)
103
104        # Update parameters and take a step using the computed gradient.
105        # The optimizer dictates the "update rule"--how the parameters are
106        # modified based on their gradients, the learning rate, etc.
107        optimizer.step()
108
109        # Update the learning rate.
110        scheduler.step()
111
112        # Calculate the average loss over all of the batches.
113        avg_train_loss = total_train_loss / len(train_dataloader)
114
115        # Measure how long this epoch took.
116        training_time = format_time(time.time() - t0)
117
118        print("")
119        print(" Average training loss: {:.2f}".format(avg_train_loss))
120        print(" Training epoch took: {}".format(training_time))
121
122        # =====
123        #             Validation
124        # =====
125        # After the completion of each training epoch, measure our performance on
126        # our validation set.
127
128        print("")
129        print("Running Validation...")
130
131        t0 = time.time()
132

```

```

133     # Put the model in evaluation mode--the dropout layers behave differently
134     # during evaluation.
135     model.eval()
136
137     # Tracking variables
138     total_eval_accuracy = 0
139     total_eval_loss = 0
140     nb_eval_steps = 0
141
142     y_pred = np.array([])
143     y_true = np.array([])
144
145     # Evaluate data for one epoch
146     for batch in validation_dataloader:
147
148         # Unpack this training batch from our dataloader.
149         #
150         # As we unpack the batch, we'll also copy each tensor to the GPU using
151         # the `to` method.
152         #
153         # `batch` contains three pytorch tensors:
154         #   [0]: input ids
155         #   [1]: attention masks
156         #   [2]: labels
157         b_input_ids = batch[0].to(device)
158         b_input_mask = batch[1].to(device)
159         b_labels = batch[2].to(device)
160
161         # Tell pytorch not to bother with constructing the compute graph during
162         # the forward pass, since this is only needed for backprop (training).
163         with torch.no_grad():
164
165             # Forward pass, calculate logit predictions.
166             # token_type_ids is the same as the "segment ids", which
167             # differentiates sentence 1 and 2 in 2-sentence tasks.
168             # The documentation for this `model` function is here:
169             # https://huggingface.co/transformers/v2.2.0/model\_doc/bert.html#transformers.Bert
170             # Get the "logits" output by the model. The "logits" are the output
171             # values prior to applying an activation function like the softmax.
172             (loss, logits) = model(b_input_ids,
173                                   token_type_ids=None,
174                                   attention_mask=b_input_mask,
175                                   labels=b_labels)
176
177         # Accumulate the validation loss.
178         total_eval_loss += loss.item()
179
180         # Move logits and labels to CPU

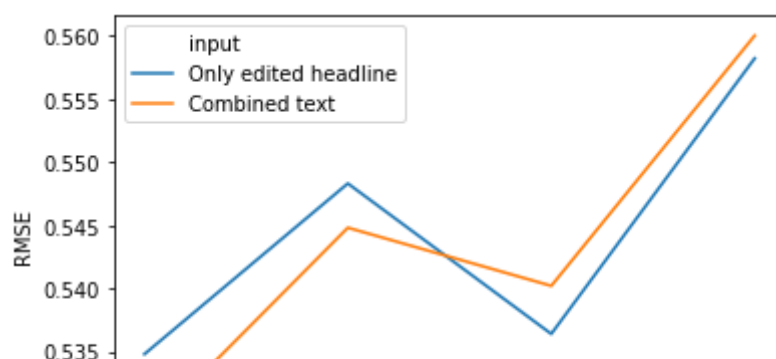
```


```

180         # move logits and labels to cpu
181         logits = logits.detach().cpu().numpy()
182         label_ids = b_labels.to('cpu').numpy()
183         y_pred = np.append(y_pred, logits)
184         y_true = np.append(y_true, label_ids)
185
186         rmse = mean_squared_error(y_true, y_pred, squared=False)
187         print(" RMSE: {0:.4f}".format(rmse))
188
189         # Calculate the average loss over all of the batches.
190         avg_val_loss = total_eval_loss / len(validation_data_loader)
191
192         # Measure how long the validation run took.
193         validation_time = format_time(time.time() - t0)
194
195         print(" Validation Loss: {0:.2f}".format(avg_val_loss))
196         print(" Validation took: {:}".format(validation_time))
197
198         # Record all statistics from this epoch.
199         training_stats.append(
200             {
201                 'epoch': epoch_i + 1,
202                 'Training Loss': avg_train_loss,
203                 'Valid. Loss': avg_val_loss,
204                 'Valid. RMSE.': rmse,
205                 'Training Time': training_time,
206                 'Validation Time': validation_time
207             }
208         )
209

```

help improve the performance.





To evaluate how good (or bad) the result is, we compare it with a model that always returns the mathematical expectation value of the mean grades based on the distribution on the training dataset. The mathematical expectation value is calculated to be approximately 0.9356, and the RMSE score of this baseline model on the validation dataset is around 0.58.

```
1  from collections import Counter
2  from sklearn.metrics import mean_squared_error
3
4  c = Counter()
5
6  for x in train_df['meanGrade']:
7      c[x] += 1
8
9  c = {k:c[k]/len(train_df['meanGrade']) for k in c}
10
11  m_e = 0
12  for x in c:
13      m_e += x * c[x]
14  # m_e = 0.9355712114933005
15
16  print(mean_squared_error(val_df['meanGrade'], [m_e]*len(val_df), squared=False))
17  # 0.5783998503042385
```

10.py hosted with ❤ by GitHub

[view raw](#)

The RMSE of around 0.53 achieved by the BERT regressor is significantly better than the baseline score of around 0.58. The BERT model seems to be able to detect some humor from given text.

However, the performance is still not so satisfactory. One reason might come from the limited range and dense distribution of the actual mean grades in the dataset. For a better model, we might need some better datasets with some larger grade range and some more uniform distribution on the grades.

## What's Next

When applying the BERT model in this article, the regressor is directly trained from the given pre-trained language model (trained with Wikipedia text). The language model can be further pre-trained with news headlines which might help improve the performance.

Also, we might try other methods to add some extra features to the BERT model. Adding features like POS of the edited words and similarities between the edited and original words might help.

[About](#) [Help](#) [Legal](#)

Get the Medium app

